

Rameesa Nadeem

Fa21-Bds-024

BIG DATA ANALYSIS PROJECT

On Textual data scrapped and ingested from website hubspot

✓ Installing Pyspark to run it on colab because gradio doesnot work on databricks

```
!apt-get install openjdk-8-jdk-headless -qq > /dev/null
!pip install pyspark
!pip install findspark
```

```
⇒ Requirement already satisfied: pyspark in /usr/local/lib/python3.11/dist-packages (3.5.4)
Requirement already satisfied: py4j==0.10.9.7 in /usr/local/lib/python3.11/dist-packages (from pyspark) (0.10.9.7)
Requirement already satisfied: findspark in /usr/local/lib/python3.11/dist-packages (2.0.1)
```

Enusuring latest version is installed

```
!pip show pyspark
```

```
⇒ Name: pyspark
Version: 3.5.4
Summary: Apache Spark Python API
Home-page: https://github.com/apache/spark/tree/master/python
Author: Spark Developers
Author-email: dev@spark.apache.org
License: http://www.apache.org/licenses/LICENSE-2.0
Location: /usr/local/lib/python3.11/dist-packages
Requires: py4j
```

Required-by:

Setting Java Path

```
import os
import findspark

os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-8-openjdk-amd64"
os.environ["SPARK_HOME"] = "/usr/local/lib/python3.11/dist-packages/pyspark" # Update this path

findspark.init()
```

Creating Spark Session

```
from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("ColabPySpark").getOrCreate()

spark
```



SparkSession - in-memory

SparkContext

[Spark UI](#)

Version

v3.5.4

Master

local[*]

AppName

ColabPySpark

▼ Scrapping and Ingesting the data from the hubspot website

```
import requests
from bs4 import BeautifulSoup
import pandas as pd

url = "https://blog.hubspot.com/marketing"
response = requests.get(url)
soup = BeautifulSoup(response.content, "html.parser")

# Extract relevant content (e.g., headlines)
data = {"content": [item.get_text(strip=True) for item in soup.find_all("h2")]}
df = pd.DataFrame(data)

# Save the data as a CSV locally
df.to_csv("marketing_data.csv", index=False)
print("Scraped data saved as 'marketing_data.csv'")
```

➡ Scraped data saved as 'marketing_data.csv'

Few Columns of csv looks like this

```
data_path = "marketing_data.csv"
data = spark.read.csv(data_path, header=True, inferSchema=True)
```

```
data.show(truncate=False)
```

➡

content
6 Steps to Create an Outstanding Marketing Plan [Free Templa...
Featured Articles
Latest articles
Artificial Intelligence
Instagram Marketing
From HubSpot's video library
Marketing Strategy
From the HubSpot Podcast Network
More content
Visit the HubSpot blogs
Subscribe to HubSpot's Newsletters
Join 600,000+ Fellow MarketersThanks for Subscribing!

Popular Features	
Free Tools	
Company	
Customers	
Partners	
+-----+	+-----+

▾
 Data Preprocessing


Removing Punctuation converting to lower case and tokenizing the text

```

from pyspark.sql.functions import col, lower, regexp_replace, split
cleaned_data = (
    data.withColumn("content", lower(col("content")))
        .withColumn("content", regexp_replace(col("content"), "[^a-zA-Z\s]", ""))
        .withColumn("tokens", split(col("content"), r"\s+"))
)

cleaned_data.show(truncate=False)

```



content	tokens
steps to create an outstanding marketing plan free template	[, steps, to, create, an, outstanding, marketing, plan, free, templa]
featured articles	[featured, articles]
latest articles	[latest, articles]
artificial intelligence	[artificial, intelligence]
instagram marketing	[instagram, marketing]
from hubspots video library	[from, hubspots, video, library]
marketing strategy	[marketing, strategy]
from the hubspot podcast network	[from, the, hubspot, podcast, network]
more content	[more, content]
visit the hubspot blogs	[visit, the, hubspot, blogs]
subscribe to hubspots newsletters	[subscribe, to, hubspots, newsletters]
join fellow marketersthanks for subscribing	[join, fellow, marketersthanks, for, subscribing]
popular features	[popular, features]
free tools	[free, tools]
company	[company]

customers	[customers]	
partners	[partners]	
+-----+	+-----+	+

Counting Words frequency

```
from pyspark.sql.functions import explode
word_freq = (
    cleaned_data.withColumn("word", explode(col("tokens")))
                .groupBy("word")
                .count()
                .orderBy("count", ascending=False)
)

word_freq.show()
```



word	count
marketing	3
articles	2
free	2
hubspots	2
the	2
from	2
hubspot	2
to	2
instagram	1
popular	1
more	1
subscribing	1
for	1
outstanding	1
library	1
content	1
featured	1
marketersthanks	1
create	1
partners	1

only showing top 20 rows

› Statistical Analysis

[] ↪ 2 cells hidden

› Machine Learning Algorithms

[] ↪ 17 cells hidden

▼ Integrating the Code with Frontend

Frontend is made using Gradio

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, lower, regexp_replace, split, explode, length
from pyspark.ml.feature import HashingTF, IDF, VectorAssembler
from pyspark.ml.clustering import KMeans
from pyspark.ml.regression import LinearRegression
from pyspark.ml.classification import LinearSVC
from pyspark.ml import Pipeline
from sklearn.cluster import DBSCAN
from sklearn.feature_extraction.text import TfidfVectorizer
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Embedding, GlobalAveragePooling1D
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
import gradio as gr
import pandas as pd
import requests
```

```

from bs4 import BeautifulSoup
import numpy as np
import gym

# Initialize Spark
spark = SparkSession.builder.appName("ColabPySpark").getOrCreate()

# Data ingestion function
def ingest_data():
    url = "https://blog.coupler.io/marketing-data-analytics/"
    response = requests.get(url)
    soup = BeautifulSoup(response.content, "html.parser")
    data = {"content": [item.get_text(strip=True) for item in soup.find_all("h2")]}
    return pd.DataFrame(data)

# Data preprocessing function
def preprocess_data():
    df = ingest_data()
    spark_df = spark.createDataFrame(df)
    cleaned_data = (
        spark_df.withColumn("content", lower(col("content")))
        .withColumn("content", regexp_replace(col("content"), "[^a-zA-Z\\s]", ""))
        .withColumn("tokens", split(col("content"), "\\s+"))
    )
    return cleaned_data

from pyspark.sql.functions import size

def compute_statistics(cleaned_data):
    word_freq = (
        cleaned_data.withColumn("word", explode(col("tokens")))
        .groupBy("word")
        .count()
    )
    variance = word_freq.selectExpr("VAR_SAMP(count) as variance").first()["variance"]

# Add content_length and word_count columns
cleaned_data = cleaned_data.withColumn("content_length", length(col("content")))
cleaned_data = cleaned_data.withColumn("word_count", size(col("tokens"))) # Use size() for array length

# Compute correlation
correlation = cleaned_data.stat.corr("content_length", "word_count")

```

```
return {"Variance": {variance}, "Correlation": {correlation}}
```

```
# K-Means Clustering
```

```
def run_kmeans(cleaned_data):  
    hashing_tf = HashingTF(inputCol="tokens", outputCol="raw_features", numFeatures=100)  
    idf = IDF(inputCol="raw_features", outputCol="features")  
    kmeans = KMeans(featuresCol="features", k=5)  
    pipeline = Pipeline(stages=[hashing_tf, idf, kmeans])  
    model = pipeline.fit(cleaned_data)  
    clusters = model.transform(cleaned_data).select("content", "prediction")  
    return clusters.toPandas()
```

```
# Linear Regression
```

```
def run_linear_regression(cleaned_data):  
    hashing_tf = HashingTF(inputCol="tokens", outputCol="raw_features", numFeatures=100)  
    idf = IDF(inputCol="raw_features", outputCol="features")  
    cleaned_data = cleaned_data.withColumn("target", (length(col("content")) % 100))  
    assembler = VectorAssembler(inputCols=["features"], outputCol="final_features")  
    lr = LinearRegression(featuresCol="final_features", labelCol="target")  
    pipeline = Pipeline(stages=[hashing_tf, idf, assembler, lr])  
    model = pipeline.fit(cleaned_data)  
    predictions = model.transform(cleaned_data).select("content", "target", "prediction")  
    return predictions.toPandas()
```

```
from pyspark.ml.classification import LogisticRegression
```

```
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
```

```
def run_classification(cleaned_data):
```

```
    # Feature extraction (TF-IDF)
```

```
    hashing_tf = HashingTF(inputCol="tokens", outputCol="raw_features", numFeatures=100)
```

```
    idf = IDF(inputCol="raw_features", outputCol="features")
```

```
    # Add mock labels for classification (e.g., 1 for even-length content, 0 for odd-length content)
```

```
    cleaned_data = cleaned_data.withColumn("label", (length(col("content")) % 2).cast("double"))
```

```
    # Train-test split
```

```
    train_data, test_data = cleaned_data.randomSplit([0.8, 0.2], seed=42)
```

```
    # Logistic Regression model
```

```
    lr = LogisticRegression(featuresCol="features", labelCol="label", maxIter=10)
```

```
    # Evaluate model
```



```

# Build pipeline
pipeline = Pipeline(stages=[hashing_tf, idf, lr])

# Train model
model = pipeline.fit(train_data)

# Make predictions
predictions = model.transform(test_data).select("content", "label", "prediction")

# Evaluate the model
evaluator = MulticlassClassificationEvaluator(
    labelCol="label", predictionCol="prediction", metricName="accuracy"
)
accuracy = evaluator.evaluate(predictions)

return predictions.toPandas(), f"Accuracy: {accuracy}"

# SVM Classification
def run_svm(cleaned_data):
    hashing_tf = HashingTF(inputCol="tokens", outputCol="raw_features", numFeatures=100)
    idf = IDF(inputCol="raw_features", outputCol="features")
    cleaned_data = cleaned_data.withColumn("label", (length(col("content")) % 2).cast("double"))
    svm = LinearSVC(featuresCol="features", labelCol="label", maxIter=10)
    pipeline = Pipeline(stages=[hashing_tf, idf, svm])
    model = pipeline.fit(cleaned_data)
    predictions = model.transform(cleaned_data).select("content", "label", "prediction")
    return predictions.toPandas()

# DBSCAN Clustering
def run_dbscan(cleaned_data):
    texts = cleaned_data.select("content").rdd.flatMap(lambda x: x).collect()
    vectorizer = TfidfVectorizer(max_features=100)
    X = vectorizer.fit_transform(texts).toarray()
    dbscan = DBSCAN(eps=0.5, min_samples=5)
    labels = dbscan.fit_predict(X)
    return pd.DataFrame({"Text": texts, "Cluster": labels})

def run_deep_learning(cleaned_data):
    texts = cleaned_data.select("content").rdd.flatMap(lambda x: x).collect()
    if not texts: # Check if the dataset is empty
        return pd.DataFrame({"Error": ["No data available for deep learning."]})

```

```

labels = [1 if len(text) % 2 == 0 else 0 for text in texts]
tokenizer = Tokenizer(num_words=5000)
tokenizer.fit_on_texts(texts)
X = pad_sequences(tokenizer.texts_to_sequences(texts), maxlen=100)

# Ensure the input data is float32
X = np.array(X, dtype="float32")
labels = np.array(labels, dtype="float32")

model = Sequential([
    Embedding(input_dim=5000, output_dim=32, input_length=100),
    GlobalAveragePooling1D(),
    Dense(16, activation='relu'),
    Dense(1, activation='sigmoid')
])
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

try:
    model.fit(X, labels, epochs=5, batch_size=32, verbose=0)
    predictions = model.predict(X)
    return pd.DataFrame({"Text": texts, "Prediction": predictions.flatten()})
except Exception as e:
    return pd.DataFrame({"Error": [str(e)]})

```

```

def run_reinforcement_learning():
    env = gym.make("CartPole-v1", new_step_api=True)
    action_space = env.action_space.n
    state_space = [20] * env.observation_space.shape[0]
    q_table = np.zeros(state_space + [action_space])
    alpha, gamma, epsilon = 0.1, 0.9, 0.1 # Learning rate, discount factor, exploration rate

```

```

def discretize_state(state, bins):
    state_bounds = list(zip(env.observation_space.low, env.observation_space.high))
    state_bounds[1] = [-3.5, 3.5] # Clip velocity
    state_bounds[3] = [-3.5, 3.5] # Clip angular velocity
    ratios = [(state[i] - state_bounds[i][0]) / (state_bounds[i][1] - state_bounds[i][0]) for i in range(len(state))]
    discrete_state = [int(r * (b - 1)) for r, b in zip(ratios, bins)]
    return tuple(np.clip(discrete_state, 0, np.array(bins) - 1))

```

5. [CartPole-v1](#) (500) - [Environment description](#) - [Function definitions](#)

```

for episode in range(500): # Increased episodes for better training
    state = discretize_state(env.reset(), state_space)
    done = False
    while not done:
        # Epsilon-greedy action selection
        if np.random.rand() < epsilon:
            action = np.random.choice(action_space) # Explore
        else:
            action = np.argmax(q_table[state]) # Exploit

        # Perform action
        next_state_raw, reward, terminated, truncated, _ = env.step(action)
        next_state = discretize_state(next_state_raw, state_space)
        done = terminated or truncated

        # Reward scaling
        reward = reward if not terminated else -100 # Penalize for termination

        # Update Q-value
        best_next_action = np.max(q_table[next_state])
        q_table[state][action] += alpha * (reward + gamma * best_next_action - q_table[state][action])

        # Move to the next state
        state = next_state

    if np.any(q_table):
        return q_table.tolist()
    else:
        return "Q-Table generation failed. Check reward structure or state discretization."

```

```

import matplotlib.pyplot as plt
import seaborn as sns
from io import BytesIO
import base64

```

```

# Visualization: Word Frequency Distribution

```

```

def plot_word_frequencies(cleaned_data):
    word_freq = (
        cleaned_data.withColumn("word", explode(col("tokens")))
        .groupBy("word")
        .count()
        .orderBy(desc("count"))
    )

```

```

        .groupBy("word")
        .count()
        .orderBy("count", ascending=False)
    ).toPandas()

plt.figure(figsize=(10, 6))
sns.barplot(x="count", y="word", data=word_freq.head(10), palette="viridis")
plt.title("Top 10 Word Frequencies")
plt.xlabel("Frequency")
plt.ylabel("Words")
return plt.gcf()

```

Visualization: K-Means Clusters

```

def plot_kmeans_clusters(kmeans_result):
    plt.figure(figsize=(10, 6))
    sns.countplot(x="prediction", data=kmeans_result, palette="viridis")
    plt.title("K-Means Cluster Distribution")
    plt.xlabel("Cluster")
    plt.ylabel("Frequency")
    return plt.gcf()

```

Visualization: DBSCAN Clusters

```

def plot_dbscan_clusters(dbscan_result):
    plt.figure(figsize=(10, 6))
    sns.countplot(x="Cluster", data=dbscan_result, palette="viridis")
    plt.title("DBSCAN Cluster Distribution")
    plt.xlabel("Cluster")
    plt.ylabel("Frequency")
    return plt.gcf()

```

```

from sklearn.decomposition import LatentDirichletAllocation

```

```

from sklearn.feature_extraction.text import CountVectorizer

```

```

def topic_modeling(cleaned_data):
    texts = cleaned_data.select("content").rdd.flatMap(lambda x: x).collect()
    vectorizer = CountVectorizer(max_features=1000, stop_words="english")
    X = vectorizer.fit_transform(texts)
    lda = LatentDirichletAllocation(n_components=5, random_state=42)
    lda.fit(X)
    topics = lda.components_
    words = vectorizer.get_feature_names_out()
    topics_words = [
        [
            word
            for word, index in zip(words, topics[i])
            if index > 0.5
        ]
        for i in range(5)
    ]
    return topics_words

```

```

topic_words = [[words[i] for i in topic.argsort()[-10:]] for topic in topics]
return pd.DataFrame({"Topic": range(1, 6), "Top Words": topic_words})

# Visualization: Reinforcement Learning Q-Table Heatmap
def plot_q_table(q_table):
    plt.figure(figsize=(12, 8))
    sns.heatmap(q_table, annot=False, cmap="viridis")
    plt.title("Reinforcement Learning Q-Table Heatmap")
    plt.xlabel("Actions")
    plt.ylabel("States")
    return plt.gcf()

# Gradio App
with gr.Blocks() as demo:
    gr.Markdown("# BIG DATA ANALYTICS PROJECT")

    # Section 1: Data Ingestion
    with gr.Row():
        with gr.Column():
            ingest_btn = gr.Button("Ingest Data")
            ingest_output = gr.Dataframe(label="Ingested Data")
            ingest_btn.click(lambda: ingest_data(), inputs=[], outputs=[ingest_output])

    # Section 2: Data Preprocessing
    with gr.Row():
        with gr.Column():
            preprocess_btn = gr.Button("Preprocess Data")
            preprocess_output = gr.Dataframe(label="Preprocessed Data")
            preprocess_btn.click(lambda: preprocess_data().toPandas(), inputs=[], outputs=[preprocess_output])

    # Section 3: Statistical Analysis
    with gr.Row():
        with gr.Column():
            stats_btn = gr.Button("Statistical Analysis")
            stats_output = gr.Textbox(label="Statistical Analysis")
            stats_btn.click(lambda: compute_statistics(preprocess_data()), inputs=[], outputs=[stats_output])

    # Section 4: Machine Learning Models
    with gr.Row():
        with gr.Column():
            kmeans_btn = gr.Button("Run K-Means Clustering")
            kmeans_output = gr.Dataframe(label="K-Means Clustering Results")
            kmeans_btn.click(lambda: kmeans(preprocess_data()), inputs=[], outputs=[kmeans_output])

```

```

kmeans_output = gr.Dataframe(label="K-Means Clusters")
kmeans_btn.click(lambda: run_kmeans(preprocess_data()), inputs=[], outputs=[kmeans_output])

with gr.Column():
    regression_btn = gr.Button("Run Linear Regression")
    regression_output = gr.Dataframe(label="Linear Regression Results")
    regression_btn.click(lambda: run_linear_regression(preprocess_data()), inputs=[], outputs=[regression_output])

with gr.Row():
    with gr.Column():
        svm_btn = gr.Button("Run SVM Classification")
        svm_output = gr.Dataframe(label="SVM Predictions")
        svm_btn.click(lambda: run_svm(preprocess_data()), inputs=[], outputs=[svm_output])

    with gr.Column():
        dbscan_btn = gr.Button("Run DBSCAN Clustering")
        dbscan_output = gr.Dataframe(label="DBSCAN Clusters")
        dbscan_btn.click(lambda: run_dbscan(preprocess_data()), inputs=[], outputs=[dbscan_output])

with gr.Row():
    with gr.Column():
        classification_btn = gr.Button("Run Classification Model")
        classification_output = gr.Dataframe(label="Classification Predictions")
        classification_accuracy = gr.Textbox(label="Classification Accuracy")
        classification_btn.click(
            lambda: run_classification(preprocess_data()),
            inputs=[],
            outputs=[classification_output, classification_accuracy]
        )

    with gr.Column():
        topic_btn = gr.Button("Run Topic Modeling")
        topic_output = gr.Dataframe(label="Topic Modeling")
        topic_btn.click(lambda: topic_modeling(preprocess_data()), inputs=[], outputs=[topic_output])

# Section 5: Deep Learning and Reinforcement Learning
with gr.Row():
    with gr.Column():
        deeplearning_btn = gr.Button("Run Deep Learning")
        deeplearning_output = gr.Dataframe(label="Deep Learning Predictions")
        deeplearning_btn.click(lambda: run_deep_learning(preprocess_data()), inputs=[], outputs=[deeplearning_output])

```

```
deeplearning_btn.click(lambda: run_deep_learning(preprocess_data()), inputs=[], outputs=[deeplearning_output])
```

```
with gr.Column():  
    rl_btn = gr.Button("Run Reinforcement Learning")  
    rl_output = gr.Textbox(label="Reinforcement Learning Q-Table")  
    rl_btn.click(run_reinforcement_learning, inputs=[], outputs=[rl_output])
```

```
# Section 6: Visualizations
```

```
with gr.Row():  
    with gr.Column():  
        word_freq_btn = gr.Button("Visualize Word Frequencies")  
        word_freq_plot = gr.Plot(label="Word Frequencies")  
        word_freq_btn.click(lambda: plot_word_frequencies(preprocess_data()), inputs=[], outputs=[word_freq_plot])  
  
    with gr.Column():  
        kmeans_vis_btn = gr.Button("Visualize K-Means Clusters")  
        kmeans_plot = gr.Plot(label="K-Means Cluster Distribution")  
        kmeans_vis_btn.click(lambda: plot_kmeans_clusters(run_kmeans(preprocess_data())), inputs=[], outputs=[kmeans_plot])  
  
    with gr.Column():  
        dbscan_vis_btn = gr.Button("Visualize DBSCAN Clusters")  
        dbscan_plot = gr.Plot(label="DBSCAN Cluster Distribution")  
        dbscan_vis_btn.click(lambda: plot_dbscan_clusters(run_dbscan(preprocess_data())), inputs=[], outputs=[dbscan_plot])  
  
    with gr.Column():  
        q_table_vis_btn = gr.Button("Visualize Q-Table")  
        q_table_plot = gr.Plot(label="Q-Table Heatmap")  
        q_table_vis_btn.click(lambda: plot_q_table(run_reinforcement_learning()), inputs=[], outputs=[q_table_plot])
```

```
demo.launch()
```



Running Gradio in a Colab notebook requires sharing enabled. Automatically setting `share=True` (you can turn this off by setting `share=False`)

Colab notebook detected. To show errors in colab notebook, set debug=True in launch()

* Running on public URL: <https://d565ed3c2bcf4db2da.gradio.live>

This share link expires in 72 hours. For free permanent hosting and GPU upgrades, run ``gradio deploy`` from the terminal in the working

BIG DATA ANALYTICS PROJECT

Ingest Data

Ingested Data

A horizontal bar with a vertical line in the center. The left side of the bar is labeled '1' and the right side is labeled '2'. There are arrows pointing outwards from the vertical line, one to the left and one to the right.

Preprocess Data

Preprocessed Data

A horizontal bar chart with two bars. The first bar is labeled '1' and the second bar is labeled '2'. The bars are of different lengths, with bar 2 being longer than bar 1. The chart has a vertical axis on the left and a horizontal axis at the bottom.

Statistical Analysis

◀ ▶